

Introduction to Ollama and LangGraph

Ollama is a tool that allows you to run large language models locally on your machine. It supports models like llama3.2, mistral, gemma, and others. By running models locally, you get 100% privacy, no API costs, and full control over your models.

LangGraph is a library for building stateful, multi-actor applications with LLMs, built on top of LangChain. It allows you to model agentic workflows as graphs, with nodes representing computational steps and edges representing routing decisions. This chatbot uses LangGraph to coordinate the document ingestion and retrieval process.